

Create database

 Reduce your Neptune Database costs by 16% with our new r8g instances (available with engine version 1.4.5 and later), delivering a better performance at a lower price than previous r6g instances. 

Instance specifications

Neptune

Amazon Neptune is a fast, reliable graph database service that makes it easy to build and run applications that work with highly connected datasets.

- High performance and scalability.
- High availability, durability, and security..
- Open graph APIs.

Type

- ☒ Provisioned
- ☐ Serverless

Engine version | [Info](#)

Neptune 1.4.8.0.R1 



► **Show filters**

Templates

Choose a template to meet your use case.

Templates

☐ **Production**
Use defaults for high availability and fast, consistent performance.

☒ **Development and testing**
This instance is intended for development use outside of a production environment.

Names

DB Cluster name

The DB cluster identifier.

fundgraph-graph

The name must be unique cross all DB clusters owned by your AWS account in the current AWS Region, case-insensitive, 1 to 60 alphanumeric characters or hyphens, first character must be a letter, it can't contain two consecutive hyphens, it can't end with a hyphen.

DB Instance name

The DB instance identifier.

fundgraph-graph-instance-1

The name must be unique cross all DB instances owned by your AWS account in the current AWS Region, case-insensitive, 1 to 60 alphanumeric characters or hyphens, first character must be a letter, it can't contain two consecutive hyphens, it can't end with a hyphen.

Cluster storage configuration - *new*

Choose the storage configuration for the Neptune cluster that best fits your application's price predictability and price performance needs.

Configuration options | [Info](#)

Database instance, storage, and I/O charges vary depending on the configuration.

☒ Neptune Standard

- Cost-effective pricing for many applications with moderate I/O usage (I/O costs <25% of total database costs).
- Pay-per-request I/O charges apply. DB instance and storage prices don't include I/O usage.

☐ Neptune I/O-Optimized

- Predictable pricing for all applications. Improved price performance for I/O-intensive applications (I/O costs >25% of total database costs).
- No additional charges for read/write I/O operations. DB instance and storage prices include I/O usage.

Instance Options

Instance class | [Info](#)

The compute and memory capacity of the DB instance.

- ☐ Memory Optimized classes (includes r and x classes)
- ☒ Burstable classes (includes t classes)

db.t3.medium

☐ Include previous generation classes

include previous generation classes

Availability and Durability

Multi-AZ deployment [Info](#)

Specifies whether the DB cluster has instances in multiple Availability Zones.

- ☐ Create read replica in different zone
- ☒ No

Network and security

Network type

To use dual-stack mode, make sure that you associate an IPv6 CIDR block with a subnet in the VPC you specify.

☒ IPv4

Your resources can communicate only over the IPv4 addressing protocol

☐ Dual-stack mode

Your resources can communicate over IPv4, IPv6, or both.

Virtual Private Cloud (VPC) [Info](#)

VPC defines the virtual networking environment for this DB instance.

Default vpc-0495c985d3fe90883 ▼



Only compatible VPCs with a corresponding DB subnet group are listed.

After a database is created, you can't change the VPC selection.

Subnet group [Info](#)

Subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.

Create new DB Subnet group ▼



Publicly accessible [Info](#)

☒ Yes

Neptune assigns a public IP address to the instances in the cluster. Amazon EC2 instances and other resources outside of the VPC can connect to your cluster, and resources inside the VPC can also connect to the cluster. Choose one or more VPC security groups to specify which resources can connect to the cluster. IAM authentication must be enabled for the cluster to allow public connectivity.

☐ No

Neptune doesn't assign a public IP address to the cluster. Only Amazon EC2 instances and other resources inside the VPC can connect to your cluster. Choose one or more VPC security groups that specify which resources can connect to the cluster.

VPC security groups

Security groups have rules authorizing connections from all the EC2 instances and devices that need to access the DB instance.

**Choose existing**

Choose existing VPC security groups

**Create new**

Create new VPC security group

Create VPC security group

Creates a new DB security group. DB security groups control access to a DB instance.

fundgraph-neptune-sg

Your data is encrypted by default with a key that AWS owns and manages for you. To choose a different key, customize your encryption settings. [Info](#)



Customize encryption settings (advanced)

► Additional settings

Tags [Info](#)

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

No tags associated with the resource.

[Add new tag](#)



Search

[Option+S]



More ▼



[Neptune](#)



[Databases](#)



Create Database

Create a Jupyter notebook to easily query your Neptune database. Jupyter notebooks are hosted and billed through SageMaker at your standard SageMaker usage rates. You can create a notebook at any time from the Neptune console.



Create notebook

Create a Neptune notebook to easily query and visualize data from your Neptune cluster.

Additional settings

DB cluster parameter group | [Info](#)

DB cluster parameter group specifies how the DB cluster is configured.

default.neptune1.4 ▼



DB instance parameter group | [Info](#)

DB instance parameter group specifies how the DB instance is configured.

default.neptune1.4 ▼



IAM DB authentication | [Info](#)



Turn on IAM Authentication

Turns on Amazon Identity and Access Management (IAM) authentication for the entire DB cluster.



Turn off IAM Authentication

Copy tags to snapshot



Turn on copy tags to snapshot

Tags are copied to any snapshot of the DB cluster that is created.



Turn off copy tags to snapshot

► **Show more**

[Cancel](#)

Create database

